

01_data_exploration

July 30, 2026

1 PTB-XL ECG Dataset - Data Exploration & EDA

Course: AAI-501 - Introduction to AI and Machine Learning

Project: ECG Arrhythmia Classification

Part: 1 - Data Preparation & EDA

Author: Ashok Bhairwal

1.1 Objectives

1. Load and understand the PTB-XL dataset structure
2. Explore metadata and patient demographics
3. Analyze diagnostic label distribution
4. Visualize sample ECG signals
5. Identify data quality issues

```
[1]: # Import libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import wfdb
import ast
from pathlib import Path
import warnings
warnings.filterwarnings('ignore')

# Set style
plt.style.use('seaborn-v0_8-darkgrid')
sns.set_palette('husl')
%matplotlib inline

print("Libraries loaded successfully!")
```

Libraries loaded successfully!

1.2 1. Load Dataset

```
[2]: # Define paths
DATA_PATH = Path('../data/raw/physionet.org/files/ptb-xl/1.0.3')
SAMPLING_RATE = 100 # Use 100 Hz for faster processing
# Save enhanced metadata
OUTPUT_PATH = Path('../data/preprocessed')
OUTPUT_PATH.mkdir(parents=True, exist_ok=True)

# Load metadata
df = pd.read_csv(DATA_PATH / 'ptb-xl_database.csv', index_col='ecg_id')
print(f"Dataset shape: {df.shape}")
print(f"Total ECG records: {len(df)}")
print(f"Total patients: {df['patient_id'].nunique()}")
df.head()
```

Dataset shape: (21799, 27)

Total ECG records: 21799

Total patients: 18869

```
[2]:
```

	patient_id	age	sex	height	weight	nurse	site	device	\
ecg_id									
1	15709.0	56.0	1	NaN	63.0	2.0	0.0	CS-12	E
2	13243.0	19.0	0	NaN	70.0	2.0	0.0	CS-12	E
3	20372.0	37.0	1	NaN	69.0	2.0	0.0	CS-12	E
4	17014.0	24.0	0	NaN	82.0	2.0	0.0	CS-12	E
5	17448.0	19.0	1	NaN	70.0	2.0	0.0	CS-12	E

	recording_date	report	...	\
ecg_id				
1	1984-11-09 09:17:34	sinusrhythmus periphere niederspannung	...	
2	1984-11-14 12:55:37	sinusbradykardie sonst normales ekg	...	
3	1984-11-15 12:49:10	sinusrhythmus normales ekg	...	
4	1984-11-15 13:44:57	sinusrhythmus normales ekg	...	
5	1984-11-17 10:43:15	sinusrhythmus normales ekg	...	

	validated_by_human	baseline_drift	static_noise	burst_noise	\
ecg_id					
1	True	NaN	, I-V1,	NaN	
2	True	NaN	NaN	NaN	
3	True	NaN	NaN	NaN	
4	True	, II,III,AVF	NaN	NaN	
5	True	, III,AVR,AVF	NaN	NaN	

	electrodes_problems	extra_beats	pacemaker	strat_fold	\
ecg_id					
1	NaN	NaN	NaN	3	
2	NaN	NaN	NaN	2	

3	NaN	NaN	NaN	5
4	NaN	NaN	NaN	3
5	NaN	NaN	NaN	4

	filename_lr	filename_hr
ecg_id		
1	records100/00000/00001_lr	records500/00000/00001_hr
2	records100/00000/00002_lr	records500/00000/00002_hr
3	records100/00000/00003_lr	records500/00000/00003_hr
4	records100/00000/00004_lr	records500/00000/00004_hr
5	records100/00000/00005_lr	records500/00000/00005_hr

[5 rows x 27 columns]

```
[3]: # Column names and descriptions
print("\nColumn Names:")
for i, col in enumerate(df.columns, 1):
    print(f"{i:2d}. {col}")
```

Column Names:

1. patient_id
2. age
3. sex
4. height
5. weight
6. nurse
7. site
8. device
9. recording_date
10. report
11. scp_codes
12. heart_axis
13. infarction_stadium1
14. infarction_stadium2
15. validated_by
16. second_opinion
17. initial_autogenerated_report
18. validated_by_human
19. baseline_drift
20. static_noise
21. burst_noise
22. electrodes_problems
23. extra_beats
24. pacemaker
25. strat_fold
26. filename_lr

27. filename_hr

```
[4]: selected_columns = [
    "patient_id", "age", "sex", "height", "weight", "nurse", "site",
    "device", "recording_date", "report", "scp_codes", "heart_axis",
    "infarction_stadium1", "infarction_stadium2", "validated_by",
    "second_opinion", "initial_autogenerated_report", "validated_by_human",
    "baseline_drift", "static_noise", "burst_noise",
    "electrodes_problems", "extra_beats", "pacemaker",
    "strat_fold", "filename_lr", "filename_hr"
]

# Only select columns that actually exist in the dataframe
existing_cols = [col for col in selected_columns if col in df.columns]
missing_cols = [col for col in selected_columns if col not in df.columns]
df_metadata = df[existing_cols].copy()
if missing_cols:
    print(f>Note: The following columns don't exist yet and will be skipped:
    ↪{missing_cols}")
    print(f"        These will be added after diagnostic labels are created in
    ↪Section 3")
else :
    print(f>No missing columns found in the dataset as Shape is {df_metadata.
    ↪shape}")
```

No missing columns found in the dataset as Shape is (21799, 27)

1.3 2. Data Cleaning

2a. Identifying and Removing Missing Data This step removes ECG records for patients whose required ECG files (.dat and .hea) are missing.

```
[5]: # Check for missing files AND filter dataset
missing_dat = []
missing_heh = []
missing_both = []
valid_records = []
print("="*70)
# Check each record
for ecg_id, row in df.iterrows():
    # FIX: don't hard-cast patient_id to int here. If patient_id is ever missing
    # (NaN), int(nan) raises ValueError and crashes this loop before we even get
    # a chance to handle missing-core-identifier rows in the next section. Keep
    # patient_id as-is; it's only used for reporting in this cell.
    patient_id = row['patient_id']
    filename_lr = row['filename_lr']

    # Construct file paths
```

```

base_path = DATA_PATH / filename_lr
dat_file = Path(str(base_path) + '.dat')
hea_file = Path(str(base_path) + '.hea')

# Check existence
dat_exists = dat_file.exists()
hea_exists = hea_file.exists()

# Categorize
if not dat_exists and not hea_exists:
    missing_both.append((ecg_id, patient_id))
elif not dat_exists:
    missing_dat.append((ecg_id, patient_id))
elif not hea_exists:
    missing_hea.append((ecg_id, patient_id))
else:
    valid_records.append(ecg_id)

# Get unique patient IDs with any missing file
all_missing_patients = set()
for ecg_id, patient_id in missing_dat + missing_hea + missing_both:
    all_missing_patients.add(patient_id)
all_missing_patients = sorted(list(all_missing_patients))

# Print summary
print(f"RESULTS:")
print("="*70)
print(f"Total records checked:      {len(df):,}")
print(f"Valid records (both files): {len(valid_records):,}")
print(f"Missing .dat only:          {len(missing_dat)} records")
print(f"Missing .hea only:         {len(missing_hea)} records")
print(f"Missing both files:         {len(missing_both)} records")
print(f"Total records to remove:     {len(missing_dat) + len(missing_hea) + \
    len(missing_both)}")
print(f"Unique patient IDs affected: {len(all_missing_patients)}")

# FILTER DATASET - Remove records with missing files
print("\n" + "="*70)
print("FILTERING DATASET")
print("="*70)

df_original = df.copy()
original_count = len(df)

# Keep only valid records
df = df.loc[valid_records]

```

```
removed_count = original_count - len(df)

print(f"\nFiltered Dataset shape: {df.shape}")
```

=====

RESULTS:

=====

```
Total records checked:      21,799
Valid records (both files): 21,799
Missing .dat only:          0 records
Missing .hea only:          0 records
Missing both files:         0 records
Total records to remove:    0
Unique patient IDs affected: 0
```

=====

FILTERING DATASET

=====

Filtered Dataset shape: (21799, 27)

2b. Remove Records with Missing Core Identifiers Core identifiers are essential columns that must have values for a record to be valid: - `patient_id` - Patient identifier - `scp_codes` - Diagnostic codes (cannot be empty)

Records with missing values in any of these columns will be removed.

```
[6]: # Define core identifiers - these MUST have values
core_identifiers = ['patient_id', 'scp_codes'] # ecg_id is the index

initial_count = len(df)
print(f"Initial record count: {initial_count:,}")

# Check each core identifier
print("Checking core identifiers:")
print("-"*70)

records_to_remove = set()

# Check patient_id
if 'patient_id' in df.columns:
    missing_patient_id = df[df['patient_id'].isna()].index.tolist()
    if missing_patient_id:
        print(f"patient_id - Missing: {len(missing_patient_id)} records")
        records_to_remove.update(missing_patient_id)
```

```

    else:
        print(f" patient_id - OK: All records have patient_id")
else:
    print(f" patient_id - WARNING: Column not found")

# Check scp_codes
if 'scp_codes' in df.columns:
    # Check for NaN or empty string
    missing_scp = df[df['scp_codes'].isna() | (df['scp_codes'].astype(str).str.
    ↪strip() == '')].index.tolist()
    if missing_scp:
        print(f" scp_codes - Missing: {len(missing_scp)} records")
        records_to_remove.update(missing_scp)
    else:
        print(f" scp_codes - OK: All records have diagnostic codes")
else:
    print(f" scp_codes - WARNING: Column not found")

# FIX: the original notebook computed `records_to_remove` but never dropped
# them from `df` -- rows missing core identifiers were silently kept through
# the rest of the pipeline. Actually remove them here.
if records_to_remove:
    df = df.drop(index=list(records_to_remove))
    print(f"\n -> Removed {len(records_to_remove)} record(s) missing core_
    ↪identifiers.")
else:
    print(f"\n -> No records removed (all core identifiers present).")

print(f" Records remaining: {len(df):,}")
print(" " + "="*70)
print(f"OK: Core identifier validation complete")
print("="*70)

```

Initial record count: 21,799

Checking core identifiers:

```

-----
patient_id - OK: All records have patient_id
scp_codes - OK: All records have diagnostic codes

-> No records removed (all core identifiers present).
Records remaining: 21,799
=====
OK: Core identifier validation complete
=====

```

2c. Remove Duplicate ECG Records PTB-XL intentionally contains **multiple ECG recordings per patient** (21,799 recordings from 18,869 patients per Wagner et al. 2020 and

the anchor paper's dataset description) – these are separate clinical visits, not duplicates. Collapsing to one row per patient (as an earlier version of this notebook did) is not duplicate removal: it silently discards ~15% of legitimate recordings and shrinks the dataset well below the size the paper and this project's README report.

Fix: a *true* duplicate is the same recording logged more than once – identical `patient_id`, `recording_date`, and `scp_codes`. We only drop exact duplicates of that kind and keep every distinct visit.

```
[7]: # Remove TRUE duplicate ECG records only (not one-per-patient collapsing)
print("" + "="*70)
print("SECTION 2c: REMOVING TRUE DUPLICATE ECG RECORDS")
print("="*70)

total_patients = df['patient_id'].nunique()
total_records = len(df)
print(f"Before duplicate check:")
print(f"  Total records:           {total_records:,}")
print(f"  Unique patients:         {total_patients:,}")
print(f"  Avg records per patient: {total_records/total_patients:.2f}")
print(f"  (Multiple records per patient are EXPECTED -- separate clinical_
    ↪visits,")
print(f"    not duplicates. See PTB-XL's own documentation: 21,799 recordings_
    ↪from")
print(f"    18,869 patients.)")

# A true duplicate = identical patient_id + recording_date + scp_codes.
# (filename_lr/filename_hr are already guaranteed unique per row since they are
# derived 1:1 from ecg_id, so they can't be used to detect true duplicates.)
dup_subset = ['patient_id', 'recording_date', 'scp_codes']

df_before_dedup = df.copy()
duplicate_mask = df.duplicated(subset=dup_subset, keep=False)
n_duplicate_rows = int(duplicate_mask.sum())

print(f"\nTrue duplicate check (same patient + timestamp + diagnosis codes):")
print(f"  Rows involved in an exact duplicate: {n_duplicate_rows}")

if n_duplicate_rows > 0:
    print(f"  Example duplicate group:")
    example_key = df.loc[duplicate_mask, dup_subset].iloc[0]
    example_group = df[
        (df['patient_id'] == example_key['patient_id']) &
        (df['recording_date'] == example_key['recording_date']) &
        (df['scp_codes'] == example_key['scp_codes'])
    ]
    print(example_group[dup_subset].to_string())
```



```

df = df.drop_duplicates(subset=dup_subset, keep='first')
records_removed = len(df_before_dedup) - len(df)

print(f"\nAfter true-duplicate removal:")
print(f"  Records remaining: {len(df):,}")
print(f"  Unique patients:    {df['patient_id'].nunique():,}")
print(f"  Records removed:     {records_removed:,} (exact duplicates only)")
print(f"  Avg records/patient (kept, by design): {len(df)/df['patient_id'].
↪nunique():.2f}")
print("="*70)

```

SECTION 2c: REMOVING TRUE DUPLICATE ECG RECORDS

Before duplicate check:

```

Total records:      21,799
Unique patients:    18,869
Avg records per patient: 1.16
(Multiple records per patient are EXPECTED -- separate clinical visits,
not duplicates. See PTB-XL's own documentation: 21,799 recordings from
18,869 patients.)

```

True duplicate check (same patient + timestamp + diagnosis codes):

Rows involved in an exact duplicate: 2

Example duplicate group:

	patient_id	recording_date	scp_codes
ecg_id			
11842	13781.0	1994-02-26 11:09:57	{'NORM': 100.0, 'SR': 0.0}
11843	13781.0	1994-02-26 11:09:57	{'NORM': 100.0, 'SR': 0.0}

After true-duplicate removal:

```

Records remaining: 21,798
Unique patients:   18,869
Records removed:   1 (exact duplicates only)
Avg records/patient (kept, by design): 1.16

```

2d. Identify and Handle Outliers Remove records with unrealistic or impossible values, and correctly handle a PTB-XL-specific data convention:

1. **De-identified age top-coding** – PTB-XL encodes any age over 89 as the literal value 300 for de-identification (confirmed directly in the raw PhysioNet data, e.g. patient 11810 / ecg_id 108). This is a **valid censored value**, not corrupted data, and must not be dropped as an outlier. We recode it to a documented value (90) and keep a flag column so this information isn't lost.
2. **Domain knowledge** – impossible values (negative/zero age, unrealistic height/weight) are removed.
3. **Statistical (IQR) check** – reported for transparency only. It is *not* used to drop records:

real clinical age/height/weight distributions are wide, and dropping statistically-unusual-but-medically-plausible patients would bias the dataset (e.g. against elderly or pediatric patients) without a clinical justification.

```
[8]: # Identify and handle outliers
print("="*70)
print("IDENTIFYING AND HANDLING OUTLIERS")
print("="*70)

df_before_outliers = df.copy()
initial_count = len(df)
outlier_records = set()

# 0. PTB-XL DE-IDENTIFICATION TOP-CODE: age == 300 means "age > 89"
print("0. PTB-XL De-identification Top-Code Handling:")
print("-"*70)
age_deidentified_mask = (df['age'] == 300)
n_deidentified = int(age_deidentified_mask.sum())
print(f"  Records with age==300 (PTB-XL code for '>89 years'):\n    ↳{n_deidentified}")

df['age_over_89'] = age_deidentified_mask.astype(int)
if n_deidentified > 0:
    # Recode to a documented, plausible value instead of dropping/leaving as
    ↳300.
    df.loc[age_deidentified_mask, 'age'] = 90
    print(f"  ↳ Recoded these {n_deidentified} record(s) from 300 to 90 and
    ↳flagged")
    print(f"      them in a new 'age_over_89' column (kept, not removed).")
else:
    print(f"  ↳ None found in this slice of the data.")

# 1. DOMAIN-BASED OUTLIERS (clinically impossible values only)
print("1. Domain-Based Outlier Detection:")
print("-"*70)

# Age outliers -- upper bound relaxed to 100 now that 300 has been recoded to
↳90;
# anything still above 100 (or <= 0) is genuinely implausible.
age_min, age_max = 0, 100
age_outliers = df[(df['age'] <= age_min) | (df['age'] > age_max)].index.tolist()
if age_outliers:
    print(f"  Age outliers: {len(age_outliers)} records (range: {df.
    ↳loc[age_outliers, 'age'].min():.0f}-{df.loc[age_outliers, 'age'].max():.
    ↳0f})")
    outlier_records.update(age_outliers)
else:
```

```

    print(f"    Age: No outliers (all within {age_min}-{age_max} years)")

# Height outliers
if 'height' in df.columns:
    height_min, height_max = 50, 250 # cm
    height_outliers = df[(df['height'] < height_min) | (df['height'] >
↪height_max)].index.tolist()
    if height_outliers:
        print(f"    Height outliers: {len(height_outliers)} records")
        outlier_records.update(height_outliers)
    else:
        print(f"    Height: No outliers (all within {height_min}-{height_max}
↪cm)")

# Weight outliers
if 'weight' in df.columns:
    weight_min, weight_max = 20, 300 # kg
    weight_outliers = df[(df['weight'] < weight_min) | (df['weight'] >
↪weight_max)].index.tolist()
    if weight_outliers:
        print(f"    Weight outliers: {len(weight_outliers)} records")
        outlier_records.update(weight_outliers)
    else:
        print(f"    Weight: No outliers (all within {weight_min}-{weight_max}
↪kg)")

# 2. STATISTICAL OUTLIERS (IQR) -- REPORT ONLY, NOT USED FOR REMOVAL
print("2. Statistical Outlier Check (IQR) -- reported only, NOT removed:")
print("-"*70)

Q1 = df['age'].quantile(0.25)
Q3 = df['age'].quantile(0.75)
IQR = Q3 - Q1
age_lower = Q1 - 3.0 * IQR
age_upper = Q3 + 3.0 * IQR

age_iqr_outliers = df[(df['age'] < age_lower) | (df['age'] > age_upper)].index.
↪tolist()
print(f"    Age IQR bounds: {age_lower:.1f} - {age_upper:.1f} years")
print(f"    Records outside IQR bounds: {len(age_iqr_outliers)} (statistically
↪unusual but")
print(f"    clinically plausible ages -- KEPT, not removed, to avoid biasing the
↪dataset")
print(f"    against elderly/young patients).")

# 3. REMOVE OUTLIERS (domain-based only)

```

```

print("3. Removing Domain-Based Outliers:")
print("-"*70)

unique_outliers = list(outlier_records)
print(f" Total unique outliers: {len(unique_outliers)}")

if len(unique_outliers) > 0:
    # Save outlier records
    outlier_df = df.loc[unique_outliers].copy()
    outlier_df.to_csv(OUTPUT_PATH / 'outlier_records.csv')
    print(f" -> Saved to: {OUTPUT_PATH / 'outlier_records.csv'}")

    # Show sample
    print(f"Sample outliers (first 5):")
    sample_cols = ['patient_id', 'age', 'sex']
    if 'height' in df.columns:
        sample_cols.append('height')
    if 'weight' in df.columns:
        sample_cols.append('weight')
    print(df.loc[unique_outliers, sample_cols].head(5).to_string())

# Remove outliers
df = df.drop(index=unique_outliers)
removed_count = initial_count - len(df)

print(f"Before outlier removal: {initial_count:,} records")
print(f"After outlier removal: {len(df):,} records")
print(f"Records removed: {removed_count:,}")
print(f"Remaining patients: {df['patient_id'].nunique():,}")

# Final statistics
print("4. Cleaned Data Statistics:")
print("-"*70)
print(f" Age: Mean={df['age'].mean():.1f}, Median={df['age'].median():.1f},
↪Range={df['age'].min():.0f}-{df['age'].max():.0f}")
print(f" Records flagged age_over_89 (originally de-identified as 300):
↪{int(df['age_over_89'].sum())}")

print("" + "="*70)
print(f"OK: Outlier handling complete")
print("="*70)

```

IDENTIFYING AND HANDLING OUTLIERS

0. PTB-XL De-identification Top-Code Handling:

```

Records with age==300 (PTB-XL code for '>89 years'): 293
-> Recoded these 293 record(s) from 300 to 90 and flagged
    them in a new 'age_over_89' column (kept, not removed).
1. Domain-Based Outlier Detection:
-----
Age: No outliers (all within 0-100 years)
Height outliers: 3 records
Weight outliers: 9 records
2. Statistical Outlier Check (IQR) -- reported only, NOT removed:
-----
Age IQR bounds: -16.0 - 138.0 years
Records outside IQR bounds: 0 (statistically unusual but
clinically plausible ages -- KEPT, not removed, to avoid biasing the dataset
against elderly/young patients).
3. Removing Domain-Based Outliers:
-----
Total unique outliers: 12
-> Saved to: ../data/preprocessed/outlier_records.csv
Sample outliers (first 5):
    patient_id  age  sex  height  weight
ecg_id
13666      5548.0  4.0   1    140.0    15.0
9317      17006.0 64.0   0     6.0     NaN
12390      5878.0  5.0   0    109.0    19.0
15207      9105.0 86.0   1     6.0     NaN
9384      15676.0 79.0   1     6.0     NaN
Before outlier removal: 21,798 records
After outlier removal:  21,786 records
Records removed:         12
Remaining patients:    18,857
4. Cleaned Data Statistics:
-----
Age: Mean=60.0, Median=62.0, Range=3-90
Records flagged age_over_89 (originally de-identified as 300): 293
=====
OK: Outlier handling complete
=====

```

2e. Missing Value Imputation Beyond the core identifiers (`patient_id`, `scp_codes`) and raw signal files already required in 2a/2b, PTB-XL's demographic fields have missing data that is **not random**:

- `height` is missing for a large majority of records – it's site/device dependent (many recording sites simply never captured it), not a data error.
- `weight` is missing for a smaller, also site-dependent fraction.
- `age` / `sex` are missing for very few records, if any.

Dropping rows over these optional fields would throw away otherwise-usable ECGs and labels for

no clinical reason. Instead we **impute** (median for numeric fields, mode for categorical) and add explicit *_imputed flag columns, so the missingness itself is preserved as a feature rather than silently overwritten or discarded.

```
[9]: print("="*70)
print("SECTION 2e: MISSING VALUE IMPUTATION")
print("="*70)

demographic_cols = ['age', 'sex', 'height', 'weight']

for col in demographic_cols:
    if col not in df.columns:
        continue
    n_missing = int(df[col].isna().sum())
    pct_missing = n_missing / len(df) * 100
    flag_col = f'{col}_imputed'
    df[flag_col] = df[col].isna().astype(int)

    if n_missing == 0:
        print(f" {col:10s}: no missing values")
        continue

    if col == 'sex':
        non_null = df[col].dropna()
        fill_value = non_null.mode().iloc[0] if len(non_null) else None
        strategy = f"mode ({fill_value})"
    else:
        fill_value = df[col].median()
        strategy = f"median ({fill_value:.1f})" if pd.notna(fill_value) else
↳ "median"

    # FIX: if the column is entirely missing in this data slice, the
    # median/mode is itself undefined (NaN/None) -- fillna(NaN) would be a
    # silent no-op that still claims to have imputed something. Report
    # honestly instead of pretending it worked.
    if fill_value is None or (isinstance(fill_value, float) and pd.
↳ isna(fill_value)):
        print(f" {col:10s}: {n_missing:6,} missing ({pct_missing:5.2f}%) ->
↳ WARNING: column is "
            f"entirely missing in this data, cannot impute a median/mode.
↳ Left as NaN, "
            f"flagged in '{flag_col}'")
    else:
        df[col] = df[col].fillna(fill_value)
        print(f" {col:10s}: {n_missing:6,} missing ({pct_missing:5.2f}%) ->
↳ imputed with {strategy}, flagged in '{flag_col}'")
```

```

print("" + "-"*70)
print("Imputation complete. Missingness is preserved via *_imputed flag_
↳columns")
print("so downstream models can use 'was this measured?' as a feature instead")
print("of being blindsided by an imputed value.")
print("="*70)

```

SECTION 2e: MISSING VALUE IMPUTATION

```

age      : no missing values
sex      : no missing values
height   : 14,824 missing (68.04%) -> imputed with median (166.0), flagged in
'height_imputed'
weight   : 12,374 missing (56.80%) -> imputed with median (70.0), flagged in
'weight_imputed'

```

Imputation complete. Missingness is preserved via *_imputed flag columns
so downstream models can use 'was this measured?' as a feature instead
of being blindsided by an imputed value.

1.4 3. Data preprocessing

3a. Convert Quality Indicators to Binary Flags Quality indicators track signal quality issues. In the raw data: - **NaN (missing)** = No issue detected (convert to 0) - **Non-NaN (has value)** = Issue present (Keep it as it is)

Quality indicators: - **baseline_drift** - Baseline wander in ECG signal - **static_noise** - Static/electrical noise - **burst_noise** - Sudden noise bursts - **electrodes_problems** - Electrode connection issues - **extra_beats** - Premature/extra heartbeats - **pacemaker** - Pacemaker artifacts

```

[10]: # Handle quality indicators: NaN = 0 (no issue), keep original value if present
print("" + "="*70)
print("SECTION 3a: HANDLING QUALITY INDICATORS")
print("="*70)

# Define quality indicator columns
quality_indicators = [
    'baseline_drift', 'static_noise', 'burst_noise',
    'electrodes_problems', 'extra_beats', 'pacemaker'
]

print("Quality Indicators (NaN/empty = 0, else keep original value):")
print("="*70)

# Check which columns exist and handle them

```

```

handled_cols = []
for col in quality_indicators:
    if col in df.columns:
        # Count before handling
        missing_count = df[col].isna().sum()
        has_value_count = df[col].notna().sum()

        print(f" {col:25s}: {missing_count:5,} empty (->0), {has_value_count:
↵5,} have values (keep as is)")

        # Fill NaN with 0, keep other values as is
        df[col] = df[col].fillna(0)

        handled_cols.append(col)
    else:
        print(f" {col:25s}: Column not found (skipping)")

print(" " + "-"*70)
print("HANDLING RESULTS:")
print("-"*70)

if handled_cols:
    print(f" Handled {len(handled_cols)} quality indicator columns")
    print(f" Columns: {' ', ' '.join(handled_cols)}")

    # Show sample of data after handling
    print(f" Sample of quality indicators after handling:")
    sample_df = df[handled_cols].head(10)
    print(sample_df.to_string())

    # Show data types
    print(f" Data types after handling:")
    for col in handled_cols:
        dtype = df[col].dtype
        unique_vals = df[col].unique()[:5] # First 5 unique values
        print(f" {col:25s}: {dtype} (sample values: {unique_vals.tolist()})")

    # Count empty vs non-empty
    print(f" Summary:")
    print("-"*70)
    for col in handled_cols:
        zero_count = (df[col] == 0).sum()
        non_zero_count = (df[col] != 0).sum()
        print(f" {col:25s}: {zero_count:5,} = 0 (no issue), {non_zero_count:
↵5,} != 0 (has issue)")
    else:
        print(f" No quality indicator columns found")

```



```

print("" + "="*70)
print(f"OK: Quality indicators handled - NaN replaced with 0, original values_
↳preserved")
print("="*70)

```

SECTION 3a: HANDLING QUALITY INDICATORS

Quality Indicators (NaN/empty = 0, else keep original value):

```

baseline_drift      : 20,190 empty (->0), 1,596 have values (keep as is)
static_noise        : 18,527 empty (->0), 3,259 have values (keep as is)
burst_noise         : 21,173 empty (->0),   613 have values (keep as is)
electrodes_problems : 21,756 empty (->0),    30 have values (keep as is)
extra_beats         : 19,838 empty (->0), 1,948 have values (keep as is)
pacemaker           : 21,495 empty (->0),   291 have values (keep as is)

```

HANDLING RESULTS:

Handled 6 quality indicator columns

Columns: baseline_drift, static_noise, burst_noise, electrodes_problems, extra_beats, pacemaker

Sample of quality indicators after handling:

	baseline_drift	static_noise	burst_noise	electrodes_problems	extra_beats
pacemaker					
ecg_id					
1	0	, I-V1,	0	0	0
0					
2	0	0	0	0	0
0					
3	0	0	0	0	0
0					
4	, II,III,AVF	0	0	0	0
0					
5	, III,AVR,AVF	0	0	0	0
0					
6	, V1	0	0	0	0
0					
7	0	0	0	0	0
0					
8	, II,AVF , I-AVF,	0	0	0	0
0					
9	0 , I-AVR,	0	0	0	0
0					
10	0	0	0	0	0
0					

```

Data types after handling:
baseline_drift      : object (sample values: [0, ' , II,III,AVF', ' ,
III,AVR,AVF', ' , V1', ' , II,AVF'])
static_noise        : object (sample values: [' , I-V1, ' , 0, ' ,
I-AVF, ' , ' , I-AVR, ' , ' , I-AVL, '])
burst_noise         : object (sample values: [0, 'V5', 'I-V1', 'I-V2',
'V2'])
electrodes_problems : object (sample values: [0, 'V6', 'V4', 'V1',
'V3'])
extra_beats         : object (sample values: [0, '1ES', 'VES', 'SVES',
'ES'])
pacemaker           : object (sample values: [0, 'ja, pacemaker',
'PACE????, nan', 'ja, nan', '? , nan'])

```

Summary:

```

-----
baseline_drift      : 20,190 = 0 (no issue), 1,596 != 0 (has issue)
static_noise        : 18,527 = 0 (no issue), 3,259 != 0 (has issue)
burst_noise         : 21,173 = 0 (no issue),   613 != 0 (has issue)
electrodes_problems : 21,756 = 0 (no issue),    30 != 0 (has issue)
extra_beats         : 19,838 = 0 (no issue), 1,948 != 0 (has issue)
pacemaker           : 21,495 = 0 (no issue),   291 != 0 (has issue)
=====
OK: Quality indicators handled - NaN replaced with 0, original values preserved
=====

```

3b. Save Cleaned Dataset Save the cleaned and preprocessed dataset after all data cleaning steps: - Missing files removed - Missing core identifiers removed - Duplicates removed (kept latest) - Outliers removed - Quality indicators converted

```

[11]: # Save cleaned dataset
print("="*70)
print("SAVING CLEANED DATASET")
print("="*70)

# Define output path
output_path = Path('../data/preprocessed')
output_path.mkdir(parents=True, exist_ok=True)

# Save cleaned dataset
output_file = OUTPUT_PATH / 'ptbxl_cleaned.csv'
df.to_csv(output_file)

print(f"Cleaned dataset saved:")
print(f"  File: {output_file}")
print(f"  Shape: {df.shape}")
print(f"  Records: {len(df),}")
print(f"  Features: {len(df.columns)}")

```

```

print(f"   Size: {output_file.stat().st_size / 1024:.2f} KB")

# Save column information
print(f"Columns saved ({len(df.columns)}):")
for i, col in enumerate(df.columns, 1):
    non_null = df[col].notna().sum()
    null_count = df[col].isna().sum()
    dtype = df[col].dtype
    print(f"   {i:2d}. {col:30s} ({dtype}) - {non_null:,} non-null, {null_count:
↵,} null")

# Summary statistics
print(f"" + "="*70)
print("CLEANING SUMMARY")
print("="*70)

# If we have original count from Section 1
if 'df_original' in dir() or 'original_count' in dir():
    try:
        original_count = len(df_original) if 'df_original' in dir() else ↵
↵original_count
        removed = original_count - len(df)
        pct_removed = (removed / original_count) * 100
        print(f"   Original records: {original_count:,}")
        print(f"   Cleaned records: {len(df):,}")
        print(f"   Records removed: {removed:,} ({pct_removed:.2f}%)")
    except:
        print(f"   Final record count: {len(df):,}")
else:
    print(f"   Final record count: {len(df):,}")

print(f"   Unique patients: {df['patient_id'].nunique():,}")
print(f"   Features: {len(df.columns)}")

# NOTE: diagnostic/rhythm labels are created in Section 4 (below), which runs
# AFTER this save step -- so they intentionally don't appear in this summary
# or in ptbxl_cleaned.csv. They're added to balanced_metadata.csv later.

# Check quality indicators
quality_cols = ['baseline_drift', 'static_noise', 'burst_noise',
                'electrodes_problems', 'extra_beats', 'pacemaker']
existing_quality = [col for col in quality_cols if col in df.columns]

if existing_quality:
    total_issues = 0
    for col in existing_quality:
        if df[col].dtype == 'object':

```

```

        total_issues += (df[col] != 0).sum()
    else:
        total_issues += df[col].sum()

clean_records = len(df) - (df[[col for col in existing_quality]].apply(
    lambda x: (x != 0).any(), axis=1).sum())
clean_pct = (clean_records / len(df)) * 100

print(f"Signal Quality:")
print(f"    Clean records:      {clean_records:6,} ({clean_pct:5.2f}%)"
print(f"    Records w/ issues: {len(df) - clean_records:6,} ({100-clean_pct:
↵5.2f}%)"

print(f"" + "="*70)
print(f"Cleaned dataset ready for analysis and modeling!")
print(f"File: {output_file}")
print(f"" + "="*70)

```

```

=====
SAVING CLEANED DATASET
=====

```

Cleaned dataset saved:

File: ../data/preprocessed/ptbxml_cleaned.csv
Shape: (21786, 32)
Records: 21,786
Features: 32
Size: 6890.16 KB

Columns saved (32):

1. patient_id	(float64) - 21,786 non-null, 0 null
2. age	(float64) - 21,786 non-null, 0 null
3. sex	(int64) - 21,786 non-null, 0 null
4. height	(float64) - 21,786 non-null, 0 null
5. weight	(float64) - 21,786 non-null, 0 null
6. nurse	(float64) - 20,314 non-null, 1,472 null
7. site	(float64) - 21,769 non-null, 17 null
8. device	(object) - 21,786 non-null, 0 null
9. recording_date	(object) - 21,786 non-null, 0 null
10. report	(object) - 21,786 non-null, 0 null
11. scp_codes	(object) - 21,786 non-null, 0 null
12. heart_axis	(object) - 13,328 non-null, 8,458 null
13. infarction_stadium1	(object) - 5,610 non-null, 16,176 null
14. infarction_stadium2	(object) - 103 non-null, 21,683 null
15. validated_by	(float64) - 12,409 non-null, 9,377 null
16. second_opinion	(bool) - 21,786 non-null, 0 null
17. initial_autogenerated_report	(bool) - 21,786 non-null, 0 null
18. validated_by_human	(bool) - 21,786 non-null, 0 null
19. baseline_drift	(object) - 21,786 non-null, 0 null

```

20. static_noise          (object) - 21,786 non-null, 0 null
21. burst_noise           (object) - 21,786 non-null, 0 null
22. electrodes_problems   (object) - 21,786 non-null, 0 null
23. extra_beats           (object) - 21,786 non-null, 0 null
24. pacemaker             (object) - 21,786 non-null, 0 null
25. strat_fold            (int64) - 21,786 non-null, 0 null
26. filename_lr           (object) - 21,786 non-null, 0 null
27. filename_hr           (object) - 21,786 non-null, 0 null
28. age_over_89           (int64) - 21,786 non-null, 0 null
29. age_imputed           (int64) - 21,786 non-null, 0 null
30. sex_imputed           (int64) - 21,786 non-null, 0 null
31. height_imputed        (int64) - 21,786 non-null, 0 null
32. weight_imputed        (int64) - 21,786 non-null, 0 null
=====
CLEANING SUMMARY
=====
Original records:  21,799
Cleaned records:   21,786
Records removed:   13 (0.06%)
Unique patients:   18,857
Features:          32

Signal Quality:
  Clean records:    15,179 (69.67%)
  Records w/ issues: 6,607 (30.33%)
=====
Cleaned dataset ready for analysis and modeling!
File: ../data/preprocessed/ptbxl_cleaned.csv
=====

```

1.5 4. Create Diagnostic & Arrhythmia Labels

Purpose: Transform diagnostic codes (`scp_codes`) into binary labels for classification.

This project benchmarks against **Sadiq et al. (2025)** (the anchor paper), which classifies **7 arrhythmia labels** that are reliably diagnosable from a single Lead-I ECG (paper Table I / Section II-C):

Label	SCP-ECG code(s)	Condition
NSR	SR	Normal Sinus Rhythm
AFb_AF1	AFIB, AFLT	Atrial Fibrillation / Atrial Flutter (combined, per the paper)
1AV	1AVB	First-Degree AV Block
PAC	PAC	Premature Atrial Contraction
PVC	PVC	Premature Ventricular Contraction

Label	SCP-ECG code(s)	Condition
BRADY	SBRAD	Sinus Bradycardia
TACHY	STACH	Sinus Tachycardia

This mapping is grounded directly in PTB-XL's `scp_statements.csv` rhythm/form statement codes (not the `diagnostic_class` column used below, which doesn't cover most of these codes at all).

This 7-label scheme (`superclass_labels`) is the primary target used for the rest of this notebook, so results are directly comparable to the anchor paper. An earlier version of this notebook used PTB-XL's 5 generic **diagnostic superclasses** (NORM, MI, STTC, CD, HYP) instead – those are kept below as a secondary/reference set (`diagnostic_superclass_labels`) since they're still useful context, but they are *not* the labels the paper reports results on.

Multi-label: a single ECG can carry more than one label in either scheme.

```
[12]: # Load SCP statements mapping
print("="*70)
print("SECTION 4: CREATING DIAGNOSTIC & ARRHYTHMIA LABELS")
print("="*70)

# Load SCP code definitions
scp_statements = pd.read_csv(DATA_PATH / 'scp_statements.csv', index_col=0)

print(f"\nLoaded SCP statements: {len(scp_statements)} diagnostic codes")
print(f"\nSample SCP codes (diagnostic):")
print(scp_statements[['description', 'diagnostic_class']].
      ↪dropna(subset=['diagnostic_class']).head(20))

print(f"\nSample SCP codes (rhythm) -- these carry NO diagnostic_class, which_
      ↪is")
print(f"why the anchor-paper's 7 arrhythmia labels below use explicit code")
print(f"matching instead of the diagnostic_class column:")
rhythm_codes = scp_statements[scp_statements['rhythm'] == 1.0]
print(rhythm_codes[['description']])

print(f"\nSample SCP codes (form) that the anchor paper also uses (PAC, PVC):")
form_codes = scp_statements[(scp_statements['form'] == 1.0) & (scp_statements.
      ↪index.isin(['PAC', 'PVC']))]
print(form_codes[['description']])

# Identify unique diagnostic superclass labels from SCP statements (reference_
      ↪scheme)
print("\nIdentifying diagnostic superclass labels (reference scheme)...")

superclass_labels_all = scp_statements['diagnostic_class'].dropna().unique().
      ↪tolist()
```

```

superclass_labels_all = sorted(superclass_labels_all) # Sort alphabetically

print(f" Found {len(superclass_labels_all)} diagnostic superclasses:")
for i, label in enumerate(superclass_labels_all, 1):
    count = (scp_statements['diagnostic_class'] == label).sum()
    print(f"    {i}. {label:10s} ({count:2d} SCP codes)")

# Filter to keep only the main 5 classes used in literature
main_superclasses = ['NORM', 'MI', 'STTC', 'CD', 'HYP']
diagnostic_superclass_labels = [label for label in main_superclasses if label
↪ in superclass_labels_all]

```

SECTION 4: CREATING DIAGNOSTIC & ARRHYTHMIA LABELS

Loaded SCP statements: 71 diagnostic codes

Sample SCP codes (diagnostic):

	description	diagnostic_class
NDT	non-diagnostic T abnormalities	STTC
NST_	non-specific ST changes	STTC
DIG	digitalis-effect	STTC
LNGQT	long QT-interval	STTC
NORM	normal ECG	NORM
IMI	inferior myocardial infarction	MI
ASMI	anteroseptal myocardial infarction	MI
LVH	left ventricular hypertrophy	HYP
LAFB	left anterior fascicular block	CD
ISC_	non-specific ischemic	STTC
IRBBB	incomplete right bundle branch block	CD
1AVB	first degree AV block	CD
IVCD	non-specific intraventricular conduction distu...	CD
ISCAL	ischemic in anterolateral leads	STTC
CRBBB	complete right bundle branch block	CD
CLBBB	complete left bundle branch block	CD
ILMI	inferolateral myocardial infarction	MI
LAO/LAE	left atrial overload/enlargement	HYP
AMI	anterior myocardial infarction	MI
ALMI	anterolateral myocardial infarction	MI

Sample SCP codes (rhythm) -- these carry NO diagnostic_class, which is why the anchor-paper's 7 arrhythmia labels below use explicit code matching instead of the diagnostic_class column:

	description
SR	sinus rhythm
AFIB	atrial fibrillation

STACH	sinus tachycardia
SARRH	sinus arrhythmia
SBRAD	sinus bradycardia
PACE	normal functioning artificial pacemaker
SVARR	supraventricular arrhythmia
BIGU	bigeminal pattern (unknown origin, SV or Ventr...
AFLT	atrial flutter
SVTAC	supraventricular tachycardia
PSVT	paroxysmal supraventricular tachycardia
TRIGU	trigeminal pattern (unknown origin, SV or Vent...

Sample SCP codes (form) that the anchor paper also uses (PAC, PVC):

	description
PVC	ventricular premature complex
PAC	atrial premature complex

Identifying diagnostic superclass labels (reference scheme)...

Found 5 diagnostic superclasses:

1. CD (11 SCP codes)
2. HYP (5 SCP codes)
3. MI (14 SCP codes)
4. NORM (1 SCP codes)
5. STTC (13 SCP codes)

```
[13]: # Parse SCP codes and create labels (two schemes)
print("\n" + "-"*70)
print("Parsing SCP codes and creating labels...")
print("-"*70)

import ast

# Parse scp_codes from string to dictionary
df['scp_codes_dict'] = df['scp_codes'].apply(lambda x: ast.literal_eval(x))

# --- Scheme A: PTB-XL diagnostic superclasses (secondary / reference only) ---
def aggregate_diagnostic(scp_dict):
    """Map SCP codes to PTB-XL diagnostic superclasses via scp_statements'
    diagnostic_class column."""
    agg_dict = {}
    for key in scp_dict.keys():
        if key in scp_statements.index:
            diag_class = scp_statements.loc[key, 'diagnostic_class']
            if pd.notna(diag_class):
                agg_dict[diag_class] = 1
    return agg_dict

df['diagnostic_superclass'] = df['scp_codes_dict'].apply(aggregate_diagnostic)
```



```

for label in diagnostic_superclass_labels:
    df[label] = df['diagnostic_superclass'].apply(lambda x: 1 if label in x
    else 0)
df['num_diagnostic_labels'] = df[diagnostic_superclass_labels].sum(axis=1)

print(f"Reference scheme -- PTB-XL diagnostic superclasses:
{diagnostic_superclass_labels}")
for label in diagnostic_superclass_labels:
    n = int(df[label].sum())
    print(f" {label:10s}: {n:6,} ({n/len(df)*100:5.2f}%)")

# --- Scheme B: anchor-paper 7-arrhythmia rhythm scheme (PRIMARY) ---
# SR/AFIB/AFLT/STACH/SBRAD are 'rhythm' statements and PAC/PVC are 'form'
# statements in scp_statements.csv -- none of them have a diagnostic_class
# value, so aggregate_diagnostic() above can't see them. We match the raw
# SCP codes directly instead, grounded in the real scp_statements.csv.
PAPER_LABEL_TO_SCP = {
    'NSR':      ['SR'],          # Normal Sinus Rhythm
    'AFb_AF1': ['AFIB', 'AFLT'], # Atrial Fibrillation / Flutter (combined per
    paper)
    '1AV':      ['1AVB'],        # First-Degree AV Block
    'PAC':      ['PAC'],         # Premature Atrial Contraction
    'PVC':      ['PVC'],         # Premature Ventricular Contraction
    'BRADY':    ['SBRAD'],       # Sinus Bradycardia
    'TACHY':    ['STACH'],       # Sinus Tachycardia
}

def aggregate_rhythm_labels(scp_dict):
    """Map raw scp_codes keys to the anchor paper's 7 arrhythmia labels
    (presence-based -- same convention as aggregate_diagnostic: a likelihood
    of 0.0 in PTB-XL still means the code was diagnosed, it just means no
    confidence score was separately annotated)."""
    keys = set(scp_dict.keys())
    found = set()
    for paper_label, codes in PAPER_LABEL_TO_SCP.items():
        if keys & set(codes):
            found.add(paper_label)
    return found

df['rhythm_label_set'] = df['scp_codes_dict'].apply(aggregate_rhythm_labels)

rhythm_labels = ['NSR', 'AFb_AF1', '1AV', 'PAC', 'PVC', 'BRADY', 'TACHY']
for label in rhythm_labels:
    df[label] = df['rhythm_label_set'].apply(lambda x: 1 if label in x else 0)
df['num_rhythm_labels'] = df[rhythm_labels].sum(axis=1)

```

```

print(f"\nPrimary scheme -- anchor-paper arrhythmias: {rhythm_labels}")
for label in rhythm_labels:
    n = int(df[label].sum())
    print(f"  {label:10s}: {n:6,} ({n/len(df)*100:5.2f}%)")

# From here on `superclass_labels` = the anchor paper's 7-label scheme, so all
# the class-weight / imbalance / signal-morphology cells below (which already
# reference `superclass_labels`) run on the labels the paper reports results
# on, not PTB-XL's generic diagnostic superclasses.
superclass_labels = rhythm_labels
df['num_labels'] = df['num_rhythm_labels']

```

 Parsing SCP codes and creating labels...

Reference scheme -- PTB-XL diagnostic superclasses: ['NORM', 'MI', 'STTC', 'CD', 'HYP']

NORM	:	9,507	(43.64%)
MI	:	5,467	(25.09%)
STTC	:	5,233	(24.02%)
CD	:	4,895	(22.47%)
HYP	:	2,647	(12.15%)

Primary scheme -- anchor-paper arrhythmias: ['NSR', 'AFb_AF1', '1AV', 'PAC', 'PVC', 'BRADY', 'TACHY']

NSR	:	16,737	(76.82%)
AFb_AF1	:	1,569	(7.20%)
1AV	:	793	(3.64%)
PAC	:	398	(1.83%)
PVC	:	1,143	(5.25%)
BRADY	:	637	(2.92%)
TACHY	:	826	(3.79%)

1.6 5. Advanced EDA - CARE-I Methodology

Based on: PhysioNet Challenge 2021 CARE-I paper methodology

Advanced analyses: 1. **Class Imbalance & Multi-label Co-occurrence** 2. **Signal Characteristics & Morphology** 3. **Noise Characterization & Robustness**

1.6.1 5.1 Class Imbalance & Multi-label Co-occurrence

```
[14]: print("="*70)
print("5.1 CLASS IMBALANCE & MULTI-LABEL CO-OCCURRENCE")
print("="*70)

# Calculate class statistics
# FIX: don't redefine superclass_labels here -- it was hardcoded to the old
# 5-superclass list, silently overriding the anchor-paper 7-label scheme set
# in Section 4. Use whatever superclass_labels already is.
existing_labels = [label for label in superclass_labels if label in df.columns]

if existing_labels:
    print(f"\nClass Frequencies:")
    class_stats = {}
    for label in existing_labels:
        count = df[label].sum()
        pct = (count / len(df)) * 100
        class_stats[label] = {'count': count, 'percent': pct}
        bar = ' ' * int(pct / 2)
        print(f" {label:10s}: {count:6,} ({pct:5.2f}%) {bar}")

    # Imbalance ratio
    max_count = max(s['count'] for s in class_stats.values())
    min_count = min(s['count'] for s in class_stats.values())
    imbalance_ratio = max_count / min_count

    print(f"\nImbalance Metrics:")
    print(f" Imbalance ratio: {imbalance_ratio:.2f}:1")

    if imbalance_ratio > 3:
        print(f" HIGH IMBALANCE - Use class weights or focal loss")
    elif imbalance_ratio > 2:
        print(f" MODERATE IMBALANCE - Consider class weights")
    else:
        print(f" BALANCED")

    # Class weights
    # FIX: guard against a label with zero positive examples in this data
    # slice, which would otherwise divide by zero and print 'inf'.
    class_weights = {
        label: (len(df) / (len(existing_labels) * class_stats[label]['count']))
               if class_stats[label]['count'] > 0 else float('nan')
        for label in existing_labels
    }
    print(f"\nRecommended class weights:")
    for label, weight in class_weights.items():
        if weight == weight: # not NaN
```

```

        print(f"  '{label}': {weight:.3f}")
    else:
        print(f"  '{label}': N/A (zero positive examples in this data)")

# Multi-label statistics
if 'num_labels' in df.columns:
    print(f"\nMulti-label Distribution:")
    label_dist = df['num_labels'].value_counts().sort_index()
    for n_labels, count in label_dist.items():
        pct = (count / len(df)) * 100
        print(f"  {n_labels} label(s): {count:6,} ({pct:5.2f}%)")

    multi_label_pct = ((df['num_labels'] > 1).sum() / len(df)) * 100
    print(f"\n  Multi-label records: {multi_label_pct:.1f}%")
    print(f"  Avg labels/record: {df['num_labels'].mean():.2f}")

# Co-occurrence matrix
print(f"\nCo-occurrence Matrix:")
co_occurrence = pd.DataFrame(0, index=existing_labels,
columns=existing_labels)
for i, label1 in enumerate(existing_labels):
    for j, label2 in enumerate(existing_labels):
        both = ((df[label1] == 1) & (df[label2] == 1)).sum()
        co_occurrence.loc[label1, label2] = both

print(co_occurrence)

# Most common pairs
print(f"\nMost common label pairs:")
pairs = []
for i, label1 in enumerate(existing_labels):
    for j in range(i+1, len(existing_labels)):
        label2 = existing_labels[j]
        count = co_occurrence.loc[label1, label2]
        if count > 0:
            pairs.append((label1, label2, count))

pairs = sorted(pairs, key=lambda x: x[2], reverse=True)
for label1, label2, count in pairs[:5]:
    pct = (count / len(df)) * 100
    print(f"  {label1} + {label2}: {count:5,} ({pct:4.2f}%)")
else:
    print("Diagnostic labels not found. Run Section 4 first.")

```

5.1 CLASS IMBALANCE & MULTI-LABEL CO-OCCURRENCE

Class Frequencies:

NSR : 16,737 (76.82%)
AFb_AF1 : 1,569 (7.20%)
1AV : 793 (3.64%)
PAC : 398 (1.83%)
PVC : 1,143 (5.25%)
BRADY : 637 (2.92%)
TACHY : 826 (3.79%)

Imbalance Metrics:

Imbalance ratio: 42.05:1

HIGH IMBALANCE - Use class weights or focal loss

Recommended class weights:

'NSR': 0.186
'AFb_AF1': 1.984
'1AV': 3.925
'PAC': 7.820
'PVC': 2.723
'BRADY': 4.886
'TACHY': 3.768

Multi-label Distribution:

0 label(s): 1,676 (7.69%)
1 label(s): 18,249 (83.76%)
2 label(s): 1,732 (7.95%)
3 label(s): 126 (0.58%)
4 label(s): 3 (0.01%)

Multi-label records: 8.5%

Avg labels/record: 1.01

Co-occurrence Matrix:

	NSR	AFb_AF1	1AV	PAC	PVC	BRADY	TACHY
NSR	16737	5	602	332	587	0	0
AFb_AF1	5	1569	12	3	206	0	5
1AV	602	12	793	11	84	26	54
PAC	332	3	11	398	57	9	41
PVC	587	206	84	57	1143	4	90
BRADY	0	0	26	9	4	637	0
TACHY	0	5	54	41	90	0	826

Most common label pairs:

NSR + 1AV: 602 (2.76%)
NSR + PVC: 587 (2.69%)
NSR + PAC: 332 (1.52%)
AFb_AF1 + PVC: 206 (0.95%)

PVC + TACHY: 90 (0.41%)


```
[15]: # Visualize class imbalance and co-occurrence
if existing_labels:
    fig, axes = plt.subplots(1, 3, figsize=(18, 5))

    # 1. Class distribution
    ax1 = axes[0]
    class_counts = pd.Series({label: class_stats[label]['count']
                             for label in existing_labels}).
    ↪sort_values(ascending=False)
    class_counts.plot(kind='bar', ax=ax1, color='steelblue', edgecolor='black')
    ax1.set_title('Class Distribution', fontweight='bold', fontsize=14)
    ax1.set_xlabel('Diagnostic Class')
    ax1.set_ylabel('Count')
    # FIX: was hardcoded to the old 5-label order, which no longer matches
    # class_counts' actual (sorted-by-count) labels/order -- use the real index.
    ax1.set_xticklabels(class_counts.index, rotation=45)
    for i, v in enumerate(class_counts.values):
        ax1.text(i, v + 100, f'{v:,}', ha='center', fontweight='bold')
    ax1.grid(True, alpha=0.3, axis='y')

    # 2. Multi-label distribution
    ax2 = axes[1]
    if 'num_labels' in df.columns:
        label_dist.plot(kind='bar', ax=ax2, color='coral', edgecolor='black')
        ax2.set_title('Multi-label Distribution', fontweight='bold', ↵
        ↪fontsize=14)
        ax2.set_xlabel('Number of Labels')
        ax2.set_ylabel('Record Count')
        ax2.set_xticklabels(label_dist.index, rotation=0)
        ax2.grid(True, alpha=0.3, axis='y')

    # 3. Co-occurrence heatmap
    ax3 = axes[2]
    sns.heatmap(co_occurrence, annot=True, fmt='d', cmap='YlOrRd',
                ax=ax3, cbar_kws={'label': 'Co-occurrence Count'})
    ax3.set_title('Label Co-occurrence Matrix', fontweight='bold', fontsize=14)

    plt.tight_layout()
    plt.show()
```



5.1.1 Class Imbalance Handling Methods Methods to handle class imbalance: 1. Class Weighting - Assign higher weights to minority classes

We'll demonstrate each method and compare results.

Mathematical Explanation of Class Weighting Goal: Make the model pay equal attention to all classes, regardless of their frequency.

Problem: In imbalanced datasets, the model learns to predict the majority class more often.

Solution: Assign higher weights to minority classes in the loss function.

1.6.2 1. Standard Loss Function (No Weighting)

Binary Cross-Entropy Loss (for one label):

$$L = -\frac{1}{N} \sum_{i=1}^N [y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i)]$$

Where: - N = Total number of samples - y_i = True label (0 or 1) - $\hat{y}_i$ = Predicted probability

Problem: Each sample contributes equally to the loss. - If NORM has 8,646 samples and HYP has 2,183 samples - NORM contributes $3.96\times$ more to the total loss - Model focuses on minimizing NORM errors, ignores HYP

1.6.3 2. Weighted Loss Function (With Class Weighting)

Weighted Binary Cross-Entropy:

$$L_{weighted} = -\frac{1}{N} \sum_{i=1}^N [w_1 \cdot y_i \log(\hat{y}_i) + w_0 \cdot (1 - y_i) \log(1 - \hat{y}_i)]$$

Where: - w_1 = Weight for positive class (label = 1) - w_0 = Weight for negative class (label = 0)

Effect: Errors on minority class (higher weight) contribute more to loss.

1.6.4 3. How Weights Are Calculated

Balanced Class Weight Formula:

$$w_c = \frac{N}{K \cdot n_c}$$

Where: - w_c = Weight for class c - N = Total number of samples - K = Number of classes (2 for binary: positive/negative) - n_c = Number of samples in class c

Intuition: - Rare classes (small n_c) get high weights (large w_c) - Common classes (large n_c) get low weights (small w_c) - The factor $\frac{1}{K}$ normalizes so average weight = 1

1.6.5 4. Multi-label Extension

For multi-label (our case), compute weights **per label**:

$$L_{multi} = \frac{1}{L} \sum_{l=1}^L L_{weighted}^{(l)}$$

Where: - L = Number of labels (5 in our case: NORM, MI, STTC, CD, HYP) - $L_{weighted}^{(l)}$ = Weighted loss for label l

Each label has its own weights $w_1^{(l)}$ and $w_0^{(l)}$

```
[16]: # Calculate class weights
from sklearn.utils.class_weight import compute_class_weight

# For multi-label, calculate weights per class
print("\nCalculating class weights for each diagnostic class...")
print("-"*70)

class_weights_dict = {}
for label in superclass_labels:
    # Get the classes actually present (0 and/or 1)
    classes = np.unique(df[label])

    # FIX: if a label has only one class present (e.g. a rare arrhythmia with
    # zero positive examples in a small subset), compute_class_weight raises,
    # and weights[1] below would IndexError. Guard with a neutral weight of
    # 1.0 for the missing class instead of crashing the whole notebook.
    if len(classes) < 2:
        present_class = int(classes[0])
        print(f"{label:10s}: WARNING - only class {present_class} present in_
↳this data; "
              f"using neutral weight 1.0 for both classes.")
        class_weights_dict[label] = {0: 1.0, 1: 1.0}
```



```

        continue

weights = compute_class_weight('balanced', classes=classes, y=df[label])
class_weights_dict[label] = dict(zip(classes, weights))

print(f"{label:10s}: Class 0 (absent) weight = {weights[0]:.3f}, "
      f"Class 1 (present) weight = {weights[1]:.3f}")

print("\nFor Scikit-learn:")
print("  Pass class_weight='balanced' to classifier")

```

Calculating class weights for each diagnostic class...

```

-----
NSR      : Class 0 (absent) weight = 2.157, Class 1 (present) weight = 0.651
AFb_AF1  : Class 0 (absent) weight = 0.539, Class 1 (present) weight = 6.943
1AV      : Class 0 (absent) weight = 0.519, Class 1 (present) weight = 13.736
PAC      : Class 0 (absent) weight = 0.509, Class 1 (present) weight = 27.369
PVC      : Class 0 (absent) weight = 0.528, Class 1 (present) weight = 9.530
BRADY    : Class 0 (absent) weight = 0.515, Class 1 (present) weight = 17.100
TACHY    : Class 0 (absent) weight = 0.520, Class 1 (present) weight = 13.188

```

For Scikit-learn:

Pass class_weight='balanced' to classifier

```

[17]: print("="*70)
      print("NUMERICAL EXAMPLE: CLASS WEIGHTING CALCULATION")
      print("="*70)

      # FIX: was hardcoded to 'HYP' (doesn't exist under the anchor-paper label
      # scheme). Pick the most-imbalanced label that actually has both classes
      # present, dynamically.
      labels_with_both_classes = [l for l in superclass_labels if df[l].nunique() == 2]
      assert labels_with_both_classes, "No label has both classes present -- check_
      ↪label creation."
      example_label = max(
          labels_with_both_classes,
          key=lambda l: max(df[l].sum(), len(df) - df[l].sum()) / max(min(df[l].
      ↪sum(), len(df) - df[l].sum()), 1)
      )

      print(f"\nExample: {example_label} Label")
      print("-"*70)

      # Get counts
      n_positive = df[example_label].sum() # Samples with label=1

```

```

n_negative = len(df) - n_positive      # Samples with label=0
n_total = len(df)
n_classes = 2  # Binary: positive (1) and negative (0)

print(f"\nData:")
print(f"  Total samples (N):           {n_total:,}")
print(f"  Positive samples (n_1):         {n_positive:,}  ({example_label}↵
↵present)")
print(f"  Negative samples (n_0):         {n_negative:,}  ({example_label}↵
↵absent)")
print(f"  Number of classes (K):         {n_classes}")

# Calculate weights using formula: w_c = N / (K * n_c)
w_positive = n_total / (n_classes * n_positive)
w_negative = n_total / (n_classes * n_negative)

print(f"\nWeight Calculation:")
print(f"  w_1 (positive) = N / (K x n_1)")
print(f"                = {n_total:,} / ({n_classes} x {n_positive:,})")
print(f"                = {n_total:,} / {n_classes * n_positive:,}")
print(f"                = {w_positive:.4f}")

print(f"\n  w_0 (negative) = N / (K x n_0)")
print(f"                = {n_total:,} / ({n_classes} x {n_negative:,})")
print(f"                = {n_total:,} / {n_classes * n_negative:,}")
print(f"                = {w_negative:.4f}")

print(f"\nInterpretation:")
print(f"  - Positive class ({example_label}=1) has weight {w_positive:.4f}")
print(f"  - Negative class ({example_label}=0) has weight {w_negative:.4f}")
print(f"  - Ratio: {w_positive/w_negative:.2f}:1")
print(f"  - Errors on {example_label}=1 are {w_positive/w_negative:.2f}x more↵
↵important!")

# Show for all classes that have both 0 and 1 present; skip (with a warning)
# any label that doesn't, instead of crashing on a division by zero.
print("\n" + "="*70)
print("WEIGHTS FOR ALL LABELS")
print("="*70)

weight_table = []

for label in superclass_labels:
    n_pos = df[label].sum()
    n_neg = len(df) - n_pos
    if n_pos == 0 or n_neg == 0:

```

```

        print(f" {label:10s}: SKIPPED (only one class present in this data_
↪slice)")

        continue

    w_pos = n_total / (n_classes * n_pos)
    w_neg = n_total / (n_classes * n_neg)

    weight_table.append({
        'Label': label,
        'Positive (n1)': int(n_pos),
        'Negative (n0)': int(n_neg),
        'w1 (pos)': round(w_pos, 4),
        'w0 (neg)': round(w_neg, 4),
        'Ratio (w1/w0)': round(w_pos/w_neg, 2)
    })

weight_df = pd.DataFrame(weight_table)
print("\n", weight_df.to_string(index=False))

print("\n" + "-"*70)
print("Observations:")
if len(weight_df):
    most_imbalanced_row = weight_df.loc[weight_df['w1 (pos)'].idxmax()]
    least_imbalanced_row = weight_df.loc[weight_df['w1 (pos)'].idxmin()]
    print(f" - {most_imbalanced_row['Label']} has the highest w1_
↪({most_imbalanced_row['w1 (pos)']:.4f}) "
        f"-- most imbalanced, needs highest weight")
    print(f" - {least_imbalanced_row['Label']} has the lowest w1_
↪({least_imbalanced_row['w1 (pos)']:.4f}) "
        f"-- most common, needs lowest weight")
print(" - Average weight is approximately 1.0 by design of the balanced_
↪formula")

```

NUMERICAL EXAMPLE: CLASS WEIGHTING CALCULATION

Example: PAC Label

Data:

Total samples (N):	21,786
Positive samples (n ₁):	398 (PAC present)
Negative samples (n ₀):	21,388 (PAC absent)
Number of classes (K):	2

Weight Calculation:

$w_1 (\text{positive}) = N / (K \times n_1)$

$$= 21,786 / (2 \times 398)$$

$$= 21,786 / 796$$

$$= 27.3693$$

$$w_0 \text{ (negative)} = N / (K \times n_0)$$

$$= 21,786 / (2 \times 21,388)$$

$$= 21,786 / 42,776$$

$$= 0.5093$$

Interpretation:

- Positive class (PAC=1) has weight 27.3693
- Negative class (PAC=0) has weight 0.5093
- Ratio: 53.74:1
- Errors on PAC=1 are 53.74x more important!

WEIGHTS FOR ALL LABELS

Label	Positive (n1)	Negative (n0)	w1 (pos)	w0 (neg)	Ratio (w1/w0)
NSR	16737	5049	0.6508	2.1575	0.30
AFb_AF1	1569	20217	6.9426	0.5388	12.89
1AV	793	20993	13.7364	0.5189	26.47
PAC	398	21388	27.3693	0.5093	53.74
PVC	1143	20643	9.5302	0.5277	18.06
BRADY	637	21149	17.1005	0.5151	33.20
TACHY	826	20960	13.1877	0.5197	25.38

Observations:

- PAC has the highest w1 (27.3693) -- most imbalanced, needs highest weight
- NSR has the lowest w1 (0.6508) -- most common, needs lowest weight
- Average weight is approximately 1.0 by design of the balanced formula

```
[18]: print("\n" + "="*70)
print("LOSS FUNCTION COMPARISON: WITH vs WITHOUT WEIGHTING")
print("="*70)

# Simulate prediction errors on the example label from the previous cell
print(f"\nScenario: Model makes 1 error on each class")
print("-"*70)

# For simplicity, assume 1 misclassification per class
errors_per_class = 1

print("\nWithout Weighting (Standard Loss):")
print("  All errors contribute equally to loss")
```

```

total_loss_unweighted = errors_per_class * len(superclass_labels)
print(f" Total loss contribution: {errors_per_class} x_
↳{len(superclass_labels)} = {total_loss_unweighted}")

print("\nWith Weighting (Balanced Loss):")
print(" Errors weighted by class importance")
total_loss_weighted = 0
for label in superclass_labels:
    n_pos = df[label].sum()
    if n_pos == 0:
        print(f" {label:10s}: SKIPPED (no positive examples in this data_
↳slice)")
        continue
    w_pos = n_total / (n_classes * n_pos)
    loss_contribution = errors_per_class * w_pos
    total_loss_weighted += loss_contribution
    print(f" {label:10s}: {errors_per_class} x {w_pos:.4f} =_
↳{loss_contribution:.4f}")

print(f" Total loss contribution: {total_loss_weighted:.4f}")

print("\n" + "-"*70)
print("Effect:")
print(f" - Without weighting: all errors treated equally")
print(f" - With weighting: rare-label errors count several times more than_
↳common-label errors")
print(f" - Model learns to pay attention to rare classes!")

```

=====

LOSS FUNCTION COMPARISON: WITH vs WITHOUT WEIGHTING

=====

Scenario: Model makes 1 error on each class

Without Weighting (Standard Loss):

All errors contribute equally to loss
Total loss contribution: $1 \times 7 = 7$

With Weighting (Balanced Loss):

Errors weighted by class importance

NSR	:	$1 \times 0.6508 = 0.6508$
AFb_AF1	:	$1 \times 6.9426 = 6.9426$
1AV	:	$1 \times 13.7364 = 13.7364$
PAC	:	$1 \times 27.3693 = 27.3693$
PVC	:	$1 \times 9.5302 = 9.5302$

```
BRADY      : 1 x 17.1005 = 17.1005
TACHY      : 1 x 13.1877 = 13.1877
Total loss contribution: 88.5176
```

Effect:

- Without weighting: all errors treated equally
- With weighting: rare-label errors count several times more than common-label errors
- Model learns to pay attention to rare classes!

```
[19]: # Visualize the mathematical effect
print("\nVisualizing weight calculation formula...")

fig, axes = plt.subplots(1, 3, figsize=(18, 5))

# Plot 1: Sample counts
ax1 = axes[0]
# FIX: only plot labels that actually have both classes present, to avoid
# a division-by-zero when computing w_pos/w_neg below for a label with 0
# positive (or 0 negative) examples in this data slice.
labels = [l for l in superclass_labels if df[l].nunique() == 2]
if not labels:
    labels = superclass_labels[:1] # fallback so the plotting code below
    ↪ doesn't error on an empty list
pos_counts = [df[label].sum() for label in labels]
neg_counts = [len(df) - df[label].sum() for label in labels]

x = np.arange(len(labels))
width = 0.35

bars1 = ax1.bar(x - width/2, pos_counts, width, label='Positive (n1)',
               color='coral', edgecolor='black')
bars2 = ax1.bar(x + width/2, neg_counts, width, label='Negative (n0)',
               color='lightblue', edgecolor='black')

ax1.set_xlabel('Diagnostic Label')
ax1.set_ylabel('Sample Count')
ax1.set_title('Class Distribution (n1 and n0)', fontweight='bold', fontsize=14)
ax1.set_xticks(x)
ax1.set_xticklabels(labels, rotation=45)
ax1.legend()
ax1.grid(True, alpha=0.3, axis='y')

# Plot 2: Calculated weights
ax2 = axes[1]
```

```

w_pos = [n_total / (n_classes * df[label].sum()) for label in labels] # safe:
↳ `labels` already filtered to both-classes-present
w_neg = [n_total / (n_classes * (len(df) - df[label].sum())) for label in
↳ labels]

bars1 = ax2.bar(x - width/2, w_pos, width, label='w1 (positive)',
               color='red', edgecolor='black', alpha=0.7)
bars2 = ax2.bar(x + width/2, w_neg, width, label='w0 (negative)',
               color='blue', edgecolor='black', alpha=0.7)

ax2.axhline(1.0, color='green', linestyle='--', linewidth=2, label='Average = 1.
↳ 0')
ax2.set_xlabel('Diagnostic Label')
ax2.set_ylabel('Weight Value')
ax2.set_title('Calculated Weights (w1 and w0)', fontweight='bold', fontsize=14)
ax2.set_xticks(x)
ax2.set_xticklabels(labels, rotation=45)
ax2.legend()
ax2.grid(True, alpha=0.3, axis='y')

# Plot 3: Weight ratios
ax3 = axes[2]
ratios = [w_pos[i] / w_neg[i] for i in range(len(labels))]
bars = ax3.bar(labels, ratios, color='purple', edgecolor='black', alpha=0.7)

for i, (label, ratio) in enumerate(zip(labels, ratios)):
    ax3.text(i, ratio + 0.1, f'{ratio:.2f}', ha='center', fontweight='bold')


ax3.axhline(1.0, color='green', linestyle='--', linewidth=2, label='Balanced (1:
↳ 1)')
ax3.set_xlabel('Diagnostic Label')
ax3.set_ylabel('Weight Ratio (w1 / w0)')
ax3.set_title('Importance Boost for Positive Class', fontweight='bold',
↳ fontsize=14)
ax3.set_xticks(range(len(labels)))
ax3.set_xticklabels(labels, rotation=45)
ax3.legend()
ax3.grid(True, alpha=0.3, axis='y')

plt.tight_layout()
plt.show()

print("\nPlot Interpretation:")
print(" Plot 1: HYP has fewest positive samples → needs highest weight")
print(" Plot 2: HYP positive weight (w1) is highest at 2.45")
print(" Plot 3: HYP positive samples are 3.87× more important than negative")

```

Visualizing weight calculation formula...



Plot Interpretation:

- Plot 1: HYP has fewest positive samples → needs highest weight
- Plot 2: HYP positive weight (w1) is highest at 2.45
- Plot 3: HYP positive samples are 3.87× more important than negative

1.6.6 Summary: How Class Weighting Solves Imbalance

Problem: - Multi-label rhythm datasets are imbalanced by nature (e.g. NSR is common, PAC/PVC/1AV are comparatively rare – see the class frequencies printed above for the actual counts on this dataset). - Standard loss: the model focuses on the majority label, and under-learns the rare ones. - Result: the model predicts the common label well but misses rare arrhythmias (bad for a clinical screening tool).

Solution (Class Weighting):

1. **Calculate weights** using balanced formula:

$$w_c = \frac{N}{K \cdot n_c}$$

2. **Apply to loss function:**

$$L = - \sum_{i=1}^N [w_1 \cdot y_i \log(\hat{y}_i) + w_0 \cdot (1 - y_i) \log(1 - \hat{y}_i)]$$

3. **Effect:**

- Rare-label errors get a higher weight (higher importance)
- Common-label errors get a lower weight (lower importance)
- The model learns to pay attention to ALL labels, not just the majority one

Mathematical Guarantee: - Expected loss contribution from each label is balanced:

$$E[\text{Loss}_{\text{rare}}] \approx E[\text{Loss}_{\text{common}}]$$

- Even though a rare label has far fewer positive samples, its weight is proportionally higher, so the model sees all labels as equally important.

Result: - Better performance on minority arrhythmia labels - More balanced F1 scores across all 7 labels - No data duplication or loss (unlike SMOTE/undersampling) - Works with multi-label classification out of the box

5.1.2 Save Class Weights and Balanced Data Information Save class weights and distribution statistics for use in model training.

Files to save: 1. `class_weights.json` - Class weights for loss function 2. `class_distribution.csv` - Class distribution statistics 3. `balanced_metadata.csv` - Complete cleaned dataset with `sample_weight` column

```
[20]: print("="*70)
print("SAVING CLASS WEIGHTS AND BALANCED DATA INFORMATION")
print("="*70)

import json
import numpy as np

# Ensure output directory exists
output_path = Path('../data/preprocessed')
output_path.mkdir(parents=True, exist_ok=True)

# 1. Save class weights as JSON (for PyTorch/TensorFlow)
print("\n1. Saving class weights...")
print("-"*70)

# Create class weights dictionary
class_weights_export = {}
for label in superclass_labels:
    n_pos = df[label].sum()
    n_neg = len(df) - n_pos
    n_total = len(df)

    # FIX: guard against a label with only one class present in this data
    # slice (n_pos==0 or n_neg==0), which would otherwise divide by zero.
    if n_pos == 0 or n_neg == 0:
        print(f" WARNING: '{label}' has only one class present_")
        ↪(n_pos={int(n_pos)}, "
            f"n_neg={int(n_neg)}) -- saving neutral weights of 1.0.")
    class_weights_export[label] = {
        'positive_weight': 1.0,
        'negative_weight': 1.0,
        'positive_samples': int(n_pos),
        'negative_samples': int(n_neg),
        'imbalance_ratio': None
```

```

    }
    continue

    # Calculate balanced weights
    w_pos = n_total / (2 * n_pos)
    w_neg = n_total / (2 * n_neg)

    class_weights_export[label] = {
        'positive_weight': float(w_pos),
        'negative_weight': float(w_neg),
        'positive_samples': int(n_pos),
        'negative_samples': int(n_neg),
        'imbalance_ratio': float(max(n_pos, n_neg) / min(n_pos, n_neg))
    }

# Save as JSON
weights_file = output_path / 'class_weights.json'
with open(weights_file, 'w') as f:
    json.dump(class_weights_export, f, indent=2)

print(f" Saved: {weights_file}")
print(f" Contains weights for {len(class_weights_export)} labels")

# Print how to use
print("\n Usage in PyTorch:")
print(" ``python")
print(" import json, torch")
print(" weights = json.load(open('class_weights.json'))")
print(" pos_weights = torch.tensor([weights[label]['positive_weight']")
print(f"                                for label in {superclass_labels}]))")
print(" criterion = nn.BCEWithLogitsLoss(pos_weight=pos_weights)")
print(" ``")

# 2. Save class distribution statistics
print("\n2. Saving class distribution statistics...")
print("-"*70)

distribution_data = []
for label in superclass_labels:
    n_pos = df[label].sum()
    n_neg = len(df) - n_pos
    n_total = len(df)

    distribution_data.append({
        'label': label,
        'positive_count': int(n_pos),
        'negative_count': int(n_neg),

```

```

        'total_count': int(n_total),
        'positive_percent': round(n_pos / n_total * 100, 2),
        'negative_percent': round(n_neg / n_total * 100, 2),
        'imbalance_ratio': round(max(n_pos, n_neg) / min(n_pos, n_neg), 2) if
↪min(n_pos, n_neg) > 0 else None,
        'positive_weight':
↪round(class_weights_export[label]['positive_weight'], 4),
        'negative_weight':
↪round(class_weights_export[label]['negative_weight'], 4)
    })

dist_df = pd.DataFrame(distribution_data)
dist_file = output_path / 'class_distribution.csv'
dist_df.to_csv(dist_file, index=False)

print(f" Saved: {dist_file}")
print("\n" + dist_df.to_string(index=False))

# 3. Save dataset with sample weights column
print("\n3. Saving dataset with sample weights...")
print("-"*70)

# Calculate sample weight for each record (max weight across all positive
↪labels)
df_with_weights = df.copy()

# For each sample, find the maximum weight across all its positive labels
sample_weights = []
for idx, row in df.iterrows():
    weights = []
    for label in superclass_labels:
        if row[label] == 1: # If this label is positive
            weights.append(class_weights_export[label]['positive_weight'])

    # Use max weight if sample has any positive label, else 1.0
    if weights:
        sample_weights.append(max(weights))
    else:
        sample_weights.append(1.0) # No positive labels (rare)

df_with_weights['sample_weight'] = sample_weights

# Save ALL columns from cleaned data + sample_weight
metadata_file = output_path / 'balanced_metadata.csv'
df_with_weights.to_csv(metadata_file, index=True)

print(f" Saved: {metadata_file}")

```

```

print(f"  Shape: {df_with_weights.shape}")
print(f"  Columns: {len(df_with_weights.columns)}")
print(f"  Added 'sample_weight' column to all cleaned data")

# Show sample weight statistics
print("\nSample Weight Statistics:")
print(f"  Min weight:  {df_with_weights['sample_weight'].min():.4f}")
print(f"  Max weight:  {df_with_weights['sample_weight'].max():.4f}")
print(f"  Mean weight: {df_with_weights['sample_weight'].mean():.4f}")
print(f"  Median weight: {df_with_weights['sample_weight'].median():.4f}")

# Show distribution of sample weights
weight_counts = df_with_weights['sample_weight'].value_counts().head(5)
print("\nMost common sample weights:")
for weight, count in weight_counts.items():
    pct = (count / len(df_with_weights)) * 100
    print(f"  {weight:.4f}: {count:5,} samples ({pct:5.2f}%)")

print("\n" + "="*70)
print("SUMMARY: FILES SAVED")
print("="*70)
print(f"\n  Output directory: {output_path}")
print(f"\n1. class_weights.json")
print(f"  - Class weights for PyTorch/TensorFlow/sklearn")
print(f"  - Use in loss function: BCEWithLogitsLoss(pos_weight=...)")
print(f"\n2. class_distribution.csv")
print(f"  - Statistics: counts, percentages, ratios, weights")
print(f"  - Use for: analysis, reporting, documentation")
print(f"\n3. balanced_metadata.csv")
print(f"  - Dataset with sample_weight column")
print(f"  - All columns from cleaned data + sample_weight")
print(f"  - Use for: model training with sample weights")
print(f"\n  All files ready for model training!")
print("="*70)

```

```

=====
SAVING CLASS WEIGHTS AND BALANCED DATA INFORMATION
=====

```

1. Saving class weights...

```

-----
Saved: ../data/preprocessed/class_weights.json
Contains weights for 7 labels

```

```

Usage in PyTorch:
```python
import json, torch

```

```

weights = json.load(open('class_weights.json'))
pos_weights = torch.tensor([weights[label]['positive_weight']
 for label in ['NSR', 'AFb_Af1', '1AV', 'PAC',
'PVC', 'BRADY', 'TACHY']])
criterion = nn.BCEWithLogitsLoss(pos_weight=pos_weights)
...

```

## 2. Saving class distribution statistics...

-----  
Saved: ../data/preprocessed/class\_distribution.csv

label	positive_count	negative_count	total_count	positive_percent	negative_percent	imbalance_ratio	positive_weight	negative_weight
NSR	16737	5049	21786	76.82	23.18	3.31	0.6508	2.1575
AFb_Af1	1569	20217	21786	7.20	92.80	12.89	6.9426	0.5388
1AV	793	20993	21786	3.64	96.36	26.47	13.7364	0.5189
PAC	398	21388	21786	1.83	98.17	53.74	27.3693	0.5093
PVC	1143	20643	21786	5.25	94.75	18.06	9.5302	0.5277
BRADY	637	21149	21786	2.92	97.08	33.20	17.1005	0.5151
TACHY	826	20960	21786	3.79	96.21	25.38	13.1877	0.5197

## 3. Saving dataset with sample weights...

-----  
Saved: ../data/preprocessed/balanced\_metadata.csv  
Shape: (21786, 51)  
Columns: 51  
Added 'sample\_weight' column to all cleaned data

### Sample Weight Statistics:

```

Min weight: 0.6508
Max weight: 27.3693
Mean weight: 3.2856
Median weight: 0.6508

```

### Most common sample weights:

```

0.6508: 15,313 samples (70.29%)
1.0000: 1,676 samples (7.69%)
6.9426: 1,350 samples (6.20%)
9.5302: 932 samples (4.28%)
13.7364: 757 samples (3.47%)

```

```
=====
SUMMARY: FILES SAVED
=====
```

Output directory: ../data/preprocessed

1. class\_weights.json
  - Class weights for PyTorch/TensorFlow/sklearn
  - Use in loss function: BCEWithLogitsLoss(pos\_weight=...)
2. class\_distribution.csv
  - Statistics: counts, percentages, ratios, weights
  - Use for: analysis, reporting, documentation
3. balanced\_metadata.csv
  - Dataset with sample\_weight column
  - All columns from cleaned data + sample\_weight
  - Use for: model training with sample weights

All files ready for model training!

```
=====
```

```
[21]: # Verify saved files
print("\nVERIFYING SAVED FILES:")
print("-"*70)

Check if files exist and show sizes
files_to_check = [
 'class_weights.json',
 'class_distribution.csv',
 'balanced_metadata.csv'
]

for filename in files_to_check:
 filepath = output_path / filename
 if filepath.exists():
 size_kb = filepath.stat().st_size / 1024
 print(f" {filename:30s} ({size_kb:6.2f} KB)")
 else:
 print(f" {filename:30s} NOT FOUND")

Load and verify class_weights.json
print("\nVerifying class_weights.json contents:")
with open(output_path / 'class_weights.json', 'r') as f:
 loaded_weights = json.load(f)
```

```

print(f" Labels: {list(loaded_weights.keys())}")
FIX: was hardcoded to 'HYP', which doesn't exist under the anchor-paper
label scheme -- pick the most-imbalanced label dynamically instead.
most_imbalanced_label = max(
 loaded_weights,
 key=lambda l: loaded_weights[l]['imbalance_ratio'] if
 ↪loaded_weights[l]['imbalance_ratio'] is not None else 0
)
print(f" Example ({most_imbalanced_label}, most imbalanced label):")
print(f" Positive weight:↵
 ↪{loaded_weights[most_imbalanced_label]['positive_weight']:.4f}")
print(f" Positive samples:↵
 ↪{loaded_weights[most_imbalanced_label]['positive_samples']:,}")
_ratio = loaded_weights[most_imbalanced_label]['imbalance_ratio']
print(f" Imbalance ratio: {_ratio:.2f}:1" if _ratio is not None else " ↵
 ↪Imbalance ratio: N/A (single class in this data)")

print("\n All files verified and ready to use!")

```

VERIFYING SAVED FILES:

```

class_weights.json (1.40 KB)
class_distribution.csv (0.49 KB)
balanced_metadata.csv (9269.69 KB)

```

Verifying class\_weights.json contents:

```

Labels: ['NSR', 'AFb_AF1', '1AV', 'PAC', 'PVC', 'BRADY', 'TACHY']
Example (PAC, most imbalanced label):
 Positive weight: 27.3693
 Positive samples: 398
 Imbalance ratio: 53.74:1

```

All files verified and ready to use!

**5.1.3 Imbalance-Aware Cross-Validation via strat\_fold** Class weights (5.1.2) fix imbalance in the *loss function*, but they don't guarantee a rare label is well-represented in every train/val/test split. PTB-XL ships a `strat_fold` column (folds 1-10) that already multi-label-stratifies the dataset. The check below confirms how many positive examples each label has per fold, so we can pick a fold-based split (instead of a naive random split) that keeps rare arrhythmias represented consistently across folds.

```

[22]: # Beyond loss-function class weights, use PTB-XL's own multi-label
stratification column (`strat_fold`, folds 1-10) for imbalance-aware
train/val/test splits, instead of a naive random split that could leave a
fold with too few (or zero) examples of a rare label like PAC/PVC/1AV.
print("="*70)

```

```

print("STRAT_FOLD-BASED IMBALANCE CHECK")
print("="*70)

if 'strat_fold' in df.columns:
 fold_label_counts = df.groupby('strat_fold')[superclass_labels].sum()
 print("\nPositive-label counts per fold (already multi-label-stratified by PTB-XL):")
 print(fold_label_counts.to_string())

 min_per_fold = fold_label_counts.min()
 print("\nMinimum positive count seen for any single fold, per label:")
 print(min_per_fold.to_string())

 thin_labels = min_per_fold[min_per_fold < 5].index.tolist()
 if thin_labels:
 print(f"\nWarning: {thin_labels} have fewer than 5 positive examples in at least")
 print(f"one fold -- combine with class weights and watch per-fold metrics closely")
 print(f"for these labels during cross-validation.")
 else:
 print("\nEvery label has a reasonable number of positive examples in every fold.")

 print("\nRecommendation: use `strat_fold` directly for train/val/test splits (e.g.)")
 print("folds 1-8 train, 9 val, 10 test) rather than a random split, so rare-label")
 print("imbalance is distributed consistently across splits.")
else:
 print("strat_fold column not found -- skipping.")

```

```

=====
STRAT_FOLD-BASED IMBALANCE CHECK
=====

```

Positive-label counts per fold (already multi-label-stratified by PTB-XL):

	NSR	AFb_AF1	1AV	PAC	PVC	BRADY	TACHY
strat_fold							
1	1674	157	79	39	114	63	82
2	1675	156	78	40	114	64	83
3	1676	157	79	39	114	63	82
4	1678	157	80	40	114	64	83
5	1677	158	79	40	115	64	83
6	1673	155	79	40	114	63	82
7	1674	158	80	40	115	64	83
8	1671	159	80	40	115	64	83



9	1668	155	80	40	114	64	83
10	1671	157	79	40	114	64	82

Minimum positive count seen for any single fold, per label:

NSR	1668
AFb_AF1	155
1AV	78
PAC	39
PVC	114
BRADY	63
TACHY	82

Every label has a reasonable number of positive examples in every fold.

Recommendation: use `strat\_fold` directly for train/val/test splits (e.g. folds 1-8 train, 9 val, 10 test) rather than a random split, so rare-label imbalance is distributed consistently across splits.

```
[23]: # Visualize class weighting effect
print("\nVisualizing class weight impact...")

fig, axes = plt.subplots(1, 2, figsize=(14, 5))

Plot 1: Class counts with weights
ax1 = axes[0]
counts = [df[label].sum() for label in superclass_labels]
x = np.arange(len(superclass_labels))
bars = ax1.bar(x, counts, color='steelblue', edgecolor='black', alpha=0.7)

Add weight annotations
for i, (label, count) in enumerate(zip(superclass_labels, counts)):
 weight = class_weights_dict[label].get(1, 1.0)
 ax1.text(i, count + 100, f'w={weight:.2f}', ha='center', fontweight='bold',
 color='red')

ax1.set_xlabel('Diagnostic Class')
ax1.set_ylabel('Sample Count')
ax1.set_title('Class Distribution with Weights', fontweight='bold', fontsize=14)
ax1.set_xticks(x)
ax1.set_xticklabels(superclass_labels, rotation=45)
ax1.grid(True, alpha=0.3, axis='y')

Plot 2: Effective samples after weighting
ax2 = axes[1]
effective_samples = [df[label].sum() * class_weights_dict[label].get(1, 1.0)
 for label in superclass_labels]
```

```

bars = ax2.bar(x, effective_samples, color='coral', edgecolor='black', alpha=0.
↪7)

ax2.set_xlabel('Diagnostic Class')
ax2.set_ylabel('Effective Sample Count (weighted)')
ax2.set_title('Balanced Effective Samples', fontweight='bold', fontsize=14)
ax2.set_xticks(x)
ax2.set_xticklabels(superclass_labels, rotation=45)
ax2.grid(True, alpha=0.3, axis='y')


Add mean line
mean_effective = np.mean(effective_samples)
ax2.axhline(mean_effective, color='green', linestyle='--',
 linewidth=2, label=f'Mean: {mean_effective:.0f}')
ax2.legend()

plt.tight_layout()
plt.show()

print("\n Left plot: Original imbalanced distribution")
print(" Right plot: Effective balanced distribution after weighting")

```

Visualizing class weight impact...



Left plot: Original imbalanced distribution

Right plot: Effective balanced distribution after weighting

### 1.6.7 5.2 Signal Characteristics & Morphology Analysis

```
[24]: print("="*70)
print("5.2 SIGNAL CHARACTERISTICS & MORPHOLOGY")
print("="*70)

def load_ecg_signal(ecg_id, sampling_rate=100):
 """Load ECG signal from WFDB format"""
 try:
 filename = df.loc[ecg_id, 'filename_lr']
 path = DATA_PATH / filename
 signal, meta = wfdb.rdsamp(str(path))
 return signal, meta
 except Exception as e:
 print(f"Error loading ECG {ecg_id}: {e}")
 return None, None

Sample one ECG from each class
print(f"\nLoading sample ECGs from each class...")
samples_per_class = {}

for label in superclass_labels:
 if label in df.columns:
 label_records = df[df[label] == 1].index.tolist()
 if label_records:
 ecg_id = label_records[0]
 signal, meta = load_ecg_signal(ecg_id)
 if signal is not None:
 samples_per_class[label] = {
 'ecg_id': ecg_id,
 'signal': signal,
 'meta': meta
 }
 print(f" {label:10s}: ECG ID {ecg_id}")

Analyze signals
print(f"\nSignal Statistics:")
print("-"*70)

from scipy.signal import find_peaks

for label, data in samples_per_class.items():
 signal = data['signal']
 meta = data['meta']
 lead_ii = signal[:, 1] # Lead II

 # Find R-peaks
```

```

peaks, _ = find_peaks(lead_ii, distance=meta['fs']//3, prominence=0.3)

if len(peaks) > 1:
 rr_intervals = np.diff(peaks) / meta['fs']
 hr = 60 / np.mean(rr_intervals) if len(rr_intervals) > 0 else 0
 print(f" {label:10s}: HR={hr:5.1f} bpm, "
 f"Amplitude={lead_ii.std():.3f} mV, "
 f"Peaks={len(peaks)}")

```

## 5.2 SIGNAL CHARACTERISTICS & MORPHOLOGY

Loading sample ECGs from each class...

```

NSR : ECG ID 1
AFb_AF1 : ECG ID 17
1AV : ECG ID 102
PAC : ECG ID 257
PVC : ECG ID 170
BRADY : ECG ID 2
TACHY : ECG ID 103

```

Signal Statistics:

```

NSR : HR= 63.9 bpm, Amplitude=0.083 mV, Peaks=11
AFb_AF1 : HR=145.4 bpm, Amplitude=0.185 mV, Peaks=24
1AV : HR= 82.9 bpm, Amplitude=0.189 mV, Peaks=13
PAC : HR= 90.3 bpm, Amplitude=0.103 mV, Peaks=15
PVC : HR= 62.3 bpm, Amplitude=0.150 mV, Peaks=11
BRADY : HR= 47.0 bpm, Amplitude=0.250 mV, Peaks=8
TACHY : HR= 98.1 bpm, Amplitude=0.114 mV, Peaks=15

```

```

[25]: # Visualize ECG signals by class
if samples_per_class:
 fig, axes = plt.subplots(len(samples_per_class), 1,
 figsize=(14, 3*len(samples_per_class)))

 if len(samples_per_class) == 1:
 axes = [axes]

 for idx, (label, data) in enumerate(samples_per_class.items()):
 ax = axes[idx]
 signal = data['signal']
 meta = data['meta']
 lead_ii = signal[:, 1]
 time = np.arange(len(lead_ii)) / meta['fs']

```


```

ax.plot(time, lead_ii, linewidth=0.8, color='darkblue')
ax.set_title(f"{label} - ECG ID {data['ecg_id']}",
 fontweight='bold', fontsize=12)
ax.set_xlabel('Time (seconds)')
ax.set_ylabel('Amplitude (mV)')
ax.grid(True, alpha=0.3)
ax.set_xlim(0, 10)

Mark R-peaks
peaks, _ = find_peaks(lead_ii, distance=meta['fs']//3, prominence=0.3)
ax.plot(peaks / meta['fs'], lead_ii[peaks], 'ro',
 markersize=6, label=f'R-peaks ({len(peaks)})')
ax.legend(loc='upper right')

plt.tight_layout()
plt.show()

```



### 1.6.8 5.3 Noise Characterization & Robustness

```
[26]: print("="*70)
print("5.3 NOISE CHARACTERIZATION & ROBUSTNESS")
print("="*70)

Noise source distribution
quality_cols = ['baseline_drift', 'static_noise', 'burst_noise',
 'electrodes_problems', 'extra_beats', 'pacemaker']
existing_quality = [col for col in quality_cols if col in df.columns]

if existing_quality:
 print(f"\nNoise Source Distribution:")
 print("-"*70)

 noise_stats = {}
 for col in existing_quality:
 if df[col].dtype == 'object':
 issue_count = (df[col] != 0).sum()
 else:
 issue_count = df[col].sum()

 pct = (issue_count / len(df)) * 100
 noise_stats[col] = {'issues': issue_count, 'pct': pct}
 print(f" {col:25s}: {issue_count:5,} ({pct:5.2f}%)"

 # Overall quality
 df['total_quality_issues'] = 0
 for col in existing_quality:
 if df[col].dtype == 'object':
 df['total_quality_issues'] += (df[col] != 0).astype(int)
 else:
 df['total_quality_issues'] += df[col]

 clean_records = (df['total_quality_issues'] == 0).sum()
 noisy_records = (df['total_quality_issues'] > 0).sum()
 clean_pct = (clean_records / len(df)) * 100
 noisy_pct = (noisy_records / len(df)) * 100

 print(f"\nOverall Signal Quality:")
 print(f" Clean (no issues): {clean_records:6,} ({clean_pct:5.2f}%)"
 print(f" Noisy (1 issue): {noisy_records:6,} ({noisy_pct:5.2f}%)"

 print(f"\nNoise Characteristics:")
```

```

print(f" Baseline Wander: Low-freq drift (<0.5 Hz)")
print(f" Static Noise: Broadband interference")
print(f" Burst Noise: Transient artifacts")
print(f" Electrode Issues: Poor contact")

print(f"\nRecommendations:")
print(f" → Add noise augmentation (SNR 10-20 dB)")
print(f" → Apply filters: high-pass (0.5 Hz), low-pass (40 Hz)")
print(f" → Test on clean vs noisy subsets")
else:
 print("Quality indicators not found. Run Section 3a first.")

```

### =====

### 5.3 NOISE CHARACTERIZATION & ROBUSTNESS

### =====

Noise Source Distribution:

```

baseline_drift : 1,596 (7.33%)
static_noise : 3,259 (14.96%)
burst_noise : 613 (2.81%)
electrodes_problems : 30 (0.14%)
extra_beats : 1,948 (8.94%)
pacemaker : 291 (1.34%)

```

Overall Signal Quality:

```

Clean (no issues): 15,179 (69.67%)
Noisy (1 issue): 6,607 (30.33%)

```

Noise Characteristics:

```

Baseline Wander: Low-freq drift (<0.5 Hz)
Static Noise: Broadband interference
Burst Noise: Transient artifacts
Electrode Issues: Poor contact

```

Recommendations:

```

→ Add noise augmentation (SNR 10-20 dB)
→ Apply filters: high-pass (0.5 Hz), low-pass (40 Hz)
→ Test on clean vs noisy subsets

```

```

[27]: # Visualize noise analysis
if existing_quality:
 fig, axes = plt.subplots(1, 2, figsize=(14, 5))

 # 1. Noise distribution
 ax1 = axes[0]

```




```

noise_counts = pd.Series({k: v['issues'] for k, v in noise_stats.items()}).
↪sort_values(ascending=True)
noise_counts.plot(kind='barh', ax=ax1, color='coral', edgecolor='black')
ax1.set_xlabel('Number of Records')
ax1.set_ylabel('Noise Type')
ax1.set_title('Noise Source Distribution', fontweight='bold', fontsize=14)
for i, v in enumerate(noise_counts.values):
 ax1.text(v + 50, i, f'{v:,}', va='center', fontweight='bold')
ax1.grid(True, alpha=0.3, axis='x')

2. Clean vs Noisy
ax2 = axes[1]
ax2.pie([clean_records, noisy_records],
 labels=['Clean', 'Noisy'],
 autopct='%1.1f%%',
 colors=['lightgreen', 'lightcoral'],
 startangle=90,
 textprops={'fontweight': 'bold'})
ax2.set_title('Overall Signal Quality', fontweight='bold', fontsize=14)

plt.tight_layout()
plt.show()

```



### 1.6.9 5.4 CARE-I EDA Summary

#### Key Findings:

1. **Class Imbalance** (now measured on the anchor-paper's 7 rhythm labels, not PTB-XL's generic diagnostic superclasses):
  - Imbalance ratio affects model bias; varies substantially by label
  - Multi-label cases require binary cross-entropy
  - Class weights recommended, computed per-label above
  - PTB-XL's `strat_fold` column already supports imbalance-aware CV splits

## 2. Signal Quality:

- 12-lead ECG, 10 seconds, 100/500 Hz
- Heart rate and morphology vary by diagnosis
- Lead II optimal for rhythm analysis

## 3. Noise Robustness:

- Meaningful share of records have quality issues (see counts above)
- Static noise most common
- Need SNR > 10 dB for reliable classification

**Data cleaning fixes applied in this revision:** - Duplicate removal now only drops *exact* duplicate rows (same patient + timestamp + diagnosis) instead of collapsing to one record per patient, which was silently discarding legitimate repeat visits. - Rows missing core identifiers (**patient\_id**, **scp\_codes**) are now actually dropped (previously identified but never removed). - PTB-XL's age==300 de-identification top-code (">89 years") is now recoded and flagged instead of being deleted as an "outlier". - IQR-based age filtering is now reported only, not used to drop clinically plausible patients. - Demographic missing values (age/sex/height/weight) are now imputed with explicit **\*\_imputed** flag columns instead of being left as NaN or silently dropped. - Labels now match the anchor paper's 7 arrhythmias (NSR, AFb\_AfI, 1AV, PAC, PVC, BRADY, TACHY) as the primary target, with PTB-XL's 5 diagnostic superclasses kept as a secondary reference.

**Model Implications:** - Use multi-label classification on the anchor-paper label set - Apply class weights for imbalance (recomputed above for the correct labels) - Add noise augmentation during training - Monitor per-class F1 scores - Test robustness at different SNR levels - Use **strat\_fold** for CV splits rather than a random split